

Database Implementation System

Project : 3

Group Member: Sathya Sai Ram Kumar(34891238) & Mohit Kalra (13906151)

Instruction to Compile and Run:

Update test.cat:

The file consists of the following variable : catalog_path, dbfile_dir, & tpch_dir paths in test.cat file.

- 1) The first line in test.cat should contain catalog_path.
- 2) The second line in test.cat should contain dbfile_dir.
- 3) The Third line in test.cat should contain tpch_dir.

Below are the instructions to compile previous assignments.

To run a1test.out:

Note: Make sure you update the variable dbfile_dir, tpch_dir and catalog_path with the appropriate location in the files a1-test.h and a1-test.cc.

```
> make clean  
> make a1test.out  
> ./a1test.out
```

To run a2test.out:

Note: Make sure you update the variable dbfile_dir, tpch_dir and catalog_path with the appropriate location in the files a2test.h and a2test.cc.

```
> make clean  
> make a2test.out  
> ./a2test.out
```

To run a2-2test.out:

Note: Make sure you update the variable dbfile_dir, tpch_dir and catalog_path with the appropriate location in the files a2-2test.h and a2-2test.cc.

```
> make clean  
> make a2-2test.out  
> ./a2-2test.out
```

To run test.out:

Note: Make sure you update the variable dbfile_dir, tpch_dir and catalog_path with the appropriate location test.h Also Make sure you have populated the dbfiles folder with the db heap files with .bin extension.

```
> make test.out  
> ./test.out
```

To run runTestCases.sh:

Note: Make sure you update the variable dbfile_dir, tpch_dir and catalog_path with the appropriate location. Also Make sure you have populated the dbfiles folder with the db heap files with .bin extension.

```
> make test.out  
> ./runTestCases.sh
```

To run gtest:

Note: Please make sure that the necessary dbfiles are already created in the dbfiles folder using a2test.out. Also, the data should be 1GB TPCB data. Following are the steps:

```
> make clean
> make gtest
> ./gtest
```

Project Flow and Structure:

Classes Created

RelationOp Class:

Base class for derivation of all the relational operators. It consists of two functions:

a. *WaitUntilDone* : This function makes sure that the currently executing thread of the relational operator is first executed before any other caller calls it.

b. *Use_n_Pages* : This function helps decides run length for BigQ objects used by several Relational Operators.

```
class RelationalOp {
    public:
        // blocks the caller until the particular relational operator
        // has run to completion
        virtual void WaitUntilDone () = 0;

        // tell us how much internal memory the operation can use
        virtual void Use_n_Pages (int n) = 0;
};
```

SelectFile Class:

SelectFile Class helps read from a dbfile and apply a given CNF predicate onto the records fetched. The records are then pushed into the output pipe.

```
class SelectFile : public RelationalOp {
    private:
        pthread_t thread;
        Record *literal;
        DBFile *inFile;
        Pipe *outPipe;
        CNF *selOp;
    public:
        void Run (DBFile &inFile, Pipe &outPipe, CNF &selOp, Record &literal);
        void WaitUntilDone ();
        void Use_n_Pages (int n);
        // why static - https://stackoverflow.com/questions/1151582/pthread-function-from-a-class
        static void* caller(void*);
        void *operation();
};
```

SelectPipe Class:

SelectPipe works in the same way as SelectFile but doesn't have a DBFile attribute attached to it. It just takes in records from the input pipe, applies the given CNF and pushes records in to the output pipe.

```
class SelectPipe : public RelationalOp {
    private:
        pthread_t thread;
        Pipe *inPipe;
        Pipe *outPipe;
        CNF *selOp;
        Record *literal;
    public:
```

```

    void Run (Pipe &inPipe, Pipe &outPipe, CNF &selOp, Record &literal);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

Project Class:

Project class is used to manipulate records of a relation in terms of the number of attributes.

```

class Project : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipe;
    Pipe *outPipe;
    int *keepMe;
    int numAttsInput;
    int numAttsOutput;
public:
    void Run (Pipe &inPipe, Pipe &outPipe, int *keepMe, int numAttsInput, int numAttsOutput);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

Join Class:

Join class implements 2 join algorithms :

- a. Sorted-Merge Join : This algorithm is used in case a proper OrderMaker object is not formed using given CNF
- b. Block-Nested Join : This algorithm is run otherwise and helps apply the CNF without using any OrderMaker objects.

```

class Join : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipeL;
    Pipe *inPipeR;
    Pipe *outPipe;
    CNF *selOp;
    Record *literal;
    int rl, mc=0, lrc=0, rrc=0;
public:
    void Run (Pipe &inPipeL, Pipe &inPipeR, Pipe &outPipe, CNF &selOp, Record &literal);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
    void MergePages(vector<Record*> lvec, Page *rp, OrderMaker &lom, OrderMaker &rom);
    void MergeRecord(Record *lr, Record *rr);
    void sortMergeJoin(Record lr, Record rr, Record m, OrderMaker &lom, OrderMaker &rom);
    void blockNestedJoin(Record lr, Record rr, Record m, OrderMaker &lom, OrderMaker &rom);
};

```

DuplicateRemoval Class:

DuplicateRemoval helps implement the **distinct** keyword of the SQL DML. It sorts the records and keep eliminating the consecutive duplicates on a given schema.

```

class DuplicateRemoval : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipe;
    Pipe *outPipe;
    Schema *mySchema;
    int rl;
public:

```

```

    void Run (Pipe &inPipe, Pipe &outPipe, Schema &mySchema);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

Sum Class:

Sum class helps implement the sum aggregation function using Function object. A function is applied on each of the valid rows of selection and the output is then pushed as a record into the output pipe.

```

class Sum : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipe;
    Pipe *outPipe;
    Function *computeMe;
public:
    void Run (Pipe &inPipe, Pipe &outPipe, Function &computeMe);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

GroupBy Class:

GroupBy works similar to sum, but applies the aggregate functions in groups of records with the same Order.

```

class GroupBy : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipe;
    Pipe *outPipe;
    OrderMaker *groupAtts;
    Function *computeMe;
    int rl;
public:
    void Run (Pipe &inPipe, Pipe &outPipe, OrderMaker &groupAtts, Function &computeMe);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

WriteOut Class:

WriteOut class helps write input pipe records into a text file as raw records. This function works in a similar way as Record.Print() but instead writes to a file than writing it out on the screen.

```

class WriteOut : public RelationalOp {
private:
    pthread_t thread;
    Pipe *inPipe;
    FILE *outFile;
    Schema *mySchema;
public:
    void Run (Pipe &inPipe, FILE *outFile, Schema &mySchema);
    void WaitUntilDone ();
    void Use_n_Pages (int n);
    static void* caller(void*);
    void *operation();
};

```

output1.txt

[illegible]

Gtest Results

```
(base) mk@mk:~/Documents/uf_docs/sem_2/Database_Implementation/workspace/DBI3$ ./gtest
[=====] Running 4 tests from 1 test suite.
[-----] Global test environment set-up.
[-----] 4 tests from QueryTesting
[ RUN      ] QueryTesting.GettingUniqueFilePath
[ OK       ] QueryTesting.GettingUniqueFilePath (1 ms)
[ RUN      ] QueryTesting.sum

** IMPORTANT: MAKE SURE THE INFORMATION BELOW IS CORRECT **
catalog location:      catalog
tpch files dir:       /home/mk/Documents/uf_docs/sem_2/Database_Implementation/10_git/tpch-dbgen/
heap files dir:       dbfiles/

[ OK       ] QueryTesting.sum (50 ms)
[ RUN      ] QueryTesting.WriteOutTesting

Number of records written to output file : 10000
[ OK       ] QueryTesting.WriteOutTesting (65 ms)
[ RUN      ] QueryTesting.DuplicateRemovalTesting
[ OK       ] QueryTesting.DuplicateRemovalTesting (46 ms)
[-----] 4 tests from QueryTesting (163 ms total)

[-----] Global test environment tear-down
[=====] 4 tests from 1 test suite ran. (163 ms total)
[ PASSED   ] 4 tests.
(base) mk@mk:~/Documents/uf_docs/sem_2/Database_Implementation/workspace/DBI3$
```

Bugs:

- 1) **In the file test.cc, query 2, line number - 133** : change clear_pipe (_p, p->schema (), true) to clear_pipe (_out, &out_sch, true);
- 2) Number of attributes for Merging and Projecting of the Record in the Join and GroupBy queries are not provided which can be passed as a parameter or can be calculated by the following formula.
Offset to first attribute- sizeof(int)/sizeof(int)
where Offset to first attribute is present in the record.
- 3) **In the file test.cc, query 6, line number - 283** : change Pipe _out (1) to Pipe _out (pipesz) as it is expecting 25 records;
- 4) **Wrong expected output given in the comments for query 2 and query 4.**
- 5) **In the file test.cc, query 6, line number - 293:** Incorrect order maker passed to group by as it should only contain the attribute "s_nationkey" and not the whole join ordermaker. Below is the code solution used which creates a new ordermaker with s_nationkey as an attribute.

```
OrderMaker grp_order;
grp_order.numAtts=1;
int n = join_sch.GetNumAtts();
Attribute *myAtts=join_sch.GetAtts();
for(int i=0;i<n;i++)
{
    if(i==3)
    {
        grp_order.whichAtts[0]=i;
        grp_order.whichTypes[0]=Int;
    }
}
```

```
G.Run (_s_ps, _out, grp_order, func);
```